

Effectful Computations with Future Conditions: A Monadic Approach to Temporal Specification

Functional Pearl

AUTHOR NAME, Institution, Country

Pre- and post-conditions are the standard vocabulary for modular program specification, yet they are silent about obligations that span multiple calls. In real Haskell code such obligations are ubiquitous: a successful TLS handshake must eventually be followed by `bye` before the context is discarded; a taken `MVar` must eventually be put back on *every* code path; a `beginTx` must eventually reach `commit` or `rollback`; a submitted job must complete with probability at least p . No single pre/post pair can express these properties, because the balancing operation lies in an unknown future context.

We present *Pledge*, a Haskell library that augments the classical contract model with *future conditions*: annotations that propagate obligations forward until they are discharged. The `Pledge` monad carries a data-dependent field `future :: a → eff`, allowing obligations to name the specific resource handle an operation returns. Monadic `bind` partially discharges the first computation's future condition against the continuation's postcondition via *subtraction*; the residual is conjoined with the continuation's own future condition.

The design is parameterised over a Composable algebra with concatenation ($\circ$), conjunction ($\wedge$), subtraction ($\setminus$), and identities `empty` and `universe`. We present four instances: trace protocols via extended regular expressions (RE, with Brzozowski-derivative subtraction and a direct LTL_f embedding); Presburger-guarded traces (GRE, discharged by $Z3$); semiring-weighted obligations (WRE, covering probabilistic and min-cost scenarios); and separation-logic heap predicates (SL, with magic wand as subtraction). All four share a single bind formula, demonstrating that trace-based, arithmetic, quantitative, and heap-ownership obligations are uniformly expressible within one algebraic interface.

CCS Concepts: • **Software and its engineering** → **Functional languages**; *Data types and structures*; *Software verification and validation*.

Additional Key Words and Phrases: temporal specification, monadic effects, regular expressions, complement, Brzozowski derivatives, Presburger arithmetic, semiring weights, separation logic, resource lifecycle, Haskell

ACM Reference Format:

Author Name. 2026. Effectful Computations with Future Conditions: A Monadic Approach to Temporal Specification: Functional Pearl. *Proc. ACM Program. Lang.* 1, HASKELL (May 2026), 25 pages.

1 Introduction

Pre- and post-conditions are the workhorse of modular program specification. A pre-condition constrains the state in which an operation may be invoked; a post-condition describes what holds immediately after it returns. Yet a great many correctness properties are not local to a single call boundary: they *span* a sequence of operations, with arbitrary computation in between.

Consider the following three scenarios from the Haskell ecosystem: i) *TLS lifecycle*. A successful handshake must eventually be followed by `bye` before the context is discarded. A post-condition on handshake records success, but cannot enforce that `bye` will be called, because the discharge

Author's Contact Information: Author Name, Institution, City, Country, author@institution.edu.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM 2475-1421/2026/5-ART

<https://doi.org/>

point lies in an unknown future context. The bracket combinator sidesteps this by lexically scoping the context, but cannot express richer message-exchange obligations or handle contexts passed across function boundaries. ii) *MVar discipline*. The `takeMVar` operation removes a value, leaving the variable empty; the caller must restore it with `putMVar` on every exit path. A post-condition records the empty state but cannot require that a future `putMVar` on the *same* variable will occur. iii) *Database savepoints*. The `transactionSave` function opens a fresh transaction that must itself be committed or rolled back; pre/post-conditions on the call cannot carry this obligation forward.

Existing remedies and their limits. The bracket function lexically scopes cleanup but cannot propagate obligations beyond the lexical scope. Both `ResourceT` [Snoyman 2011] and `Acquire/Managed` thread cleanup actions through every `bind` but model only the cleanup half of an obligation, discharging it uniformly at scope exit rather than by application-specific operations. Linear Haskell [Bernardy et al. 2018] enforces “consume exactly once” but cannot express branching discharge obligations or quantitative commitments. Parameterised and indexed monads [Atkey 2009] encode protocol states at the type level but require `RebindableSyntax` or `QualifiedDo`, produce opaque type errors, and are confined to a single domain per design.

Our approach. We propose the `Pledge` monad as a lightweight, compositional mechanism for specifying such obligations, requiring no language extensions beyond standard Haskell type classes. A value of type `Pledge eff a` carries three annotations: a precondition `pre :: eff`, a postcondition `post :: eff`, and a data-dependent future condition `future :: a → eff`. Because `future` ranges over the return value, an obligation can name the actual resource handle produced at runtime: `tlsHandshake` can declare that *whichever* `Context` it returns must eventually receive `bye`. Monadic `bind` partially discharges the first computation’s future via subtraction against the continuation’s postcondition; the residual is conjoined with the continuation’s own future condition. The user annotates individual operations only; the monad propagates obligations automatically.

The `Composable` typeclass abstracts the specification algebra into five operations: concatenation ($\diamond$), conjunction ($\wedge$), subtraction ($\setminus$), and their respective identities `empty` and `universe`. It admits four independent instances:

- (*Trace protocol*, RE) Extended regular expressions over events, with complement as a first-class constructor. Subtraction is implemented via Brzowski derivatives; LTL_f operators embed directly as RE terms without automaton construction. This instance handles `takeMVar/putMVar` discipline, TLS handshake/bye pairing, transaction lifecycle, and lock protocols.
- (*Arithmetic bounds*, GRE) Each RE is paired with a Presburger arithmetic predicate over program variables, checked by an SMT solver (Z3 via SBV). A single annotation simultaneously enforces a protocol ordering *and* a numeric bound, catching, for example, counter overflow before any execution takes place.
- (*Quantitative obligation*, WRE) Boolean membership is generalised to a semiring weight. The probability semiring captures obligations of the form “this job completes with probability at least p ”; the tropical semiring gives minimum-cost scheduling obligations. A weight of `szero` in the residual signals an unfulfilled quantitative commitment.
- (*Heap ownership*, SL) Separation-logic heap predicates, with separating conjunction as sequential composition and magic wand as subtraction. Future conditions describe what heap state must hold when the computation completes; a leaked `Cell` in the residual `post` names the unreleased ownership precisely.

Contributions.

- (1) We define the Pledge monad (§3), parameterised over a Composable algebra, with $\text{pre} :: \text{eff}$, $\text{post} :: \text{eff}$, and a data-dependent future $a \rightarrow \text{eff}$ field, and verify that the monad laws hold for the $(\text{ret}, \text{post}, \text{future})$ triple (§9).
- (2) We instantiate the algebra to extended regular expressions over events with complement, using Brzozowski derivatives to implement subtraction (§2.1), and show that LTL_f operators embed directly as RE terms (§2.3).
- (3) We extend the RE instance to a guarded form (GRE, §5), pairing each RE with a Presburger arithmetic predicate over program variables, evaluated by an SMT solver (Z3 via SBV) at membership-check time.
- (4) We provide a weighted RE instance (WRE, §6), parameterised over a semiring, that generalises Boolean membership to a quantitative weight, recovering probabilistic and min-cost future conditions as special cases.
- (5) We provide a fourth, independent instantiation to separation-logic heap predicates (§7), demonstrating that the Composable interface extends beyond trace-based specifications.
- (6) We demonstrate the approach on ten use cases across all four instances (§8), showing that resource leaks, ordering violations, arithmetic overflows, probabilistic obligations, and heap-ownership errors all surface as non-trivial residual annotations.

2 Background

2.1 Extended Regular Expressions over Events

Events are the atomic units of observation; each carries a name and a list of typed arguments.

```

data Term = Str String | Num Int | List [Term]
data Event = Atom String [Term]  -- concrete event, e.g. Atom "send" [Num 1]
           | Wildcard             -- pattern matching any single event
data RE = Bot | Epsilon | Single Event
         | Seq RE RE | Or RE RE | And RE RE | Star RE | Not RE

```

The specification language is an extended regular expression type that adds complement as a first-class constructor:

$$r ::= \emptyset \mid \varepsilon \mid e \mid r_1 \cdot r_2 \mid r_1 \vee r_2 \mid r_1 \wedge r_2 \mid r^* \mid \neg r$$

The constructors are standard. The empty language $\emptyset$ (**Bot**) accepts no trace. The empty-trace language ε (**Epsilon**) accepts only the empty sequence. A single-event term e (**Single** Event) matches exactly one event; **Wildcard** is a pattern that matches any single event. Concatenation $r_1 \cdot r_2$ (**Seq**) accepts a trace if it splits into a prefix in $L(r_1)$ followed by a suffix in $L(r_2)$. Union $r_1 \vee r_2$ (**Or**) accepts traces in either language. The Kleene star r^* (**Star**) accepts zero or more repetitions of traces from $L(r)$; Σ^* is the set of all finite event sequences over the alphabet Σ (the universal language). Intersection $r_1 \wedge r_2$ (**And**) is derived as $\neg(\neg r_1 \vee \neg r_2)$ via De Morgan. Complement $\neg r$ (**Not**) is the set of all traces in Σ^* not in $L(r)$. Moreover, we use the following five derived forms throughout:

$\text{anything} \triangleq \neg \emptyset$	(the universal language Σ^*)
$\text{finally}(e) \triangleq \text{anything} \cdot e \cdot \text{anything}$	(e must occur eventually)
$\text{globally}(e) \triangleq e^*$	(e at every step)
$\text{never}(e) \triangleq \neg \text{finally}(e)$	(e must never occur)
$\text{noUntil}(e, g) \triangleq \neg((\Sigma \setminus \{g\})^* \cdot e \cdot \text{anything})$	(e must not occur before g)

$\text{noUntil}(e, g)$ is nullable (the empty continuation satisfies it) and its derivative with respect to g is Σ^* , meaning once g occurs the restriction on e is lifted. This is the key combinator for *conditional* safety: e is forbidden until g re-enables it.

2.2 Algebraic Complement and Brzowski Derivatives

The Brzowski derivative [Brzowski 1964] $\partial_a(r)$ computes the language of continuations after reading event a . The key law for complement is:

$$\partial_a(\neg r) = \neg(\partial_a(r)) \quad (1)$$

The nullability function $\nu(r)$ ($= \text{True}$ iff $\varepsilon \in L(r)$) satisfies $\nu(\neg r) = \neg \nu(r)$. Together, equations (1) and the De Morgan laws

$$\neg(r_1 \vee r_2) = \neg r_1 \wedge \neg r_2 \quad \neg(r_1 \wedge r_2) = \neg r_1 \vee \neg r_2$$

allow complement to be propagated and simplified *purely algebraically*, without building a DFA. The normaliser additionally applies $\neg \neg r = r$, $\neg \emptyset = \neg \emptyset$, and $\neg \neg \emptyset = \emptyset$.

In Haskell:

```
nullable :: RE → Bool
nullable (Not r)      = not (nullable r)      -- nu(Not r) = not(nu(r))
nullable (Star _)     = True
nullable (Seq r1 r2)  = nullable r1 && nullable r2
nullable (Or r1 r2)   = nullable r1 || nullable r2
nullable (And r1 r2)  = nullable r1 && nullable r2
nullable Epsilon      = True
nullable _            = False

derivative :: Event → RE → RE
derivative e (Not r)   = Not (derivative e r)  -- d_a(Not r) = Not(d_a(r))
derivative e (Single p) = if matchEvent e p then Epsilon else Bot
derivative e (Seq r1 r2)
  | nullable r1 = Or (Seq (derivative e r1) r2) (derivative e r2)
  | otherwise   = Seq (derivative e r1) r2
derivative e (Or r1 r2) = Or (derivative e r1) (derivative e r2)
derivative e (And r1 r2) = And (derivative e r1) (derivative e r2)
derivative e (Star r)   = Seq (derivative e r) (Star r)
derivative _ _         = Bot
```

The `firstWith` function returns the set of events (drawn from a supplied finite alphabet Σ_0) that can begin a word in $L(r)$, driving the subtraction algorithm. The complement case requires particular care. Naively writing `first (Not r) = first r` is *unsound*: $L(\neg r)$ can begin with events *not* in $\text{first}(r)$, so the naive rule explores the wrong direction. The correct characterisation is:

$$e \in \text{first}(\neg r, \Sigma_0) \iff e \in \Sigma_0 \wedge [\partial_e(r)] \neq \neg \emptyset \quad (2)$$

That is, e can open a word in $L(\neg r)$ iff after consuming e , the residual $\partial_e(r)$ does *not* cover the whole future; there exists some continuation that stays outside $L(r)$. The alphabet Σ_0 is collected from an RE via atoms, which recursively gathers all concrete (Atom) events, excluding Wildcard.

```
atoms :: RE → [Event]
atoms (Single Wildcard) = []
```

```

atoms (Single e)      = [e]
atoms (Seq r1 r2)     = atoms r1 `union` atoms r2
atoms (Or r1 r2)      = atoms r1 `union` atoms r2
atoms (And r1 r2)     = atoms r1 `union` atoms r2
atoms (Star r)        = atoms r
atoms (Not r)         = atoms r
atoms _               = []

firstWith :: [Event] → RE → [Event]
firstWith _ Bot       = []
firstWith _ Epsilon   = []
firstWith _ (Single e) = [e]
firstWith alph (Seq r1 r2)
  | nullable r1       = firstWith alph r1 `union` firstWith alph r2
  | otherwise         = firstWith alph r1
firstWith alph (Or r1 r2) = firstWith alph r1 `union` firstWith alph r2
firstWith alph (And r1 r2) = [e | e <- firstWith alph r1,
                                   e `elem` firstWith alph r2]
firstWith alph (Star r)  = firstWith alph r
firstWith alph (Not r)   = [e | e <- alph,
                               not (isTotal (normalize (derivative e r)))]
  where isTotal (Not Bot) = True; isTotal _ = False

```

The function `first r = firstWith (atoms r) r` is a convenient wrapper for standalone use; inside subtraction both operands contribute to the alphabet (see §4).

Illustrative cases. Table 1 gives representative evaluations of `firstWith` with **Not**. The first four rows confirm that complement does not exclude any alphabet event when the derivative is non-total. The double-negation row shows correct restoration of the original first set: $\partial_b(-a) = \Sigma^*$ is total, so b is filtered out. In $\neg(a \cdot \Sigma^*)$, $\partial_a(a \cdot \Sigma^*)$ normalises to Σ^* , excluding a ; only b survives. Similarly, $\neg(a \cdot \Sigma^* \vee b \cdot \Sigma^*)$ excludes both a and b , leaving only c , while the De Morgan row verifies that $\neg(\neg a \wedge \neg b) = a \vee b$ yields $\{a, b\}$.

Table 1. Selected `firstWith` evaluations with **Not**. Σ_0 is the explicitly supplied alphabet. Rows marked $\dagger$ produce a *strict* subset of Σ_0 .

RE r	Σ_0	<code>firstWith</code> Σ_0 ($\neg r$)	Rationale
a	$\{a, b\}$	$\{a, b\}$	$\partial_a(a) = \varepsilon$, $\partial_b(a) = \emptyset$; neither Σ^*
$a \cdot b$	$\{a, b, c\}$	$\{a, b, c\}$	all derivatives non- Σ^*
$a \vee b$	$\{a, b, c\}$	$\{a, b, c\}$	$\partial_c(a \vee b) = \emptyset$, still non- Σ^*
a^*	$\{a, b\}$	$\{a, b\}$	$\partial_a(a^*) = a^*$, non- Σ^* ; b starts $[b]$
$\neg a$ (double neg.)	$\{a, b\}$	$[a]$	$\partial_b(\neg a) = \Sigma^* \Rightarrow b$ excluded †
$a \cdot \Sigma^*$	$\{a, b\}$	$[b]$	$\partial_a(a \cdot \Sigma^*) = \Sigma^* \Rightarrow a$ excluded †
$a \cdot \Sigma^* \vee b \cdot \Sigma^*$	$\{a, b, c\}$	$[c]$	both a, b excluded; only c avoids †
$\neg a \wedge \neg b$	$\{a, b, c\}$	$\{a, b\}$	De Morgan: $\neg(\neg a \vee \neg b)$; ∂_c is $\Sigma^{*\dagger}$

2.3 Embedding LTL_f

Because complement is a first-class constructor, the standard LTL_f operators translate directly into RE, with no automaton construction needed. Table 2 gives the full mapping.

Table 2. LTL_f to RE translation. $\Sigma^1 \triangleq \text{Single Wildcard}$, $\neg\emptyset \triangleq \neg\emptyset$.

LTL formula	RE $\llbracket \cdot \rrbracket$
$\top$	$\neg\emptyset$
$\perp$	$\emptyset$
e	$\{e\}$
$\neg\phi$	$\neg\llbracket \phi \rrbracket$
$\phi \wedge \psi$	$\llbracket \phi \rrbracket \cap \llbracket \psi \rrbracket$
$\phi \vee \psi$	$\llbracket \phi \rrbracket \cup \llbracket \psi \rrbracket$
$X\phi$	$\Sigma^1 \cdot \llbracket \phi \rrbracket$
$F\phi$	$\neg\emptyset \cdot \llbracket \phi \rrbracket$
$G\phi$	$\neg(\neg\emptyset \cdot \neg\llbracket \phi \rrbracket)$
$\phi U \psi$	$\text{step}(\phi)^* \cdot \llbracket \psi \rrbracket$ (ϕ propositional)

The complement operator does the heavy lifting: $G\phi \triangleq \neg F\neg\phi$ unfolds to $\neg(\neg\emptyset \cdot \neg\llbracket \phi \rrbracket)$ using two applications of $\neg$; $F\phi$ and $X\phi$ need only concatenation. The U case uses $\text{step}(\phi)$, the length-1 RE for a single event satisfying ϕ , which is defined only when ϕ is propositional.

3 The Pledge Monad

3.1 The Composable Class

We abstract the specification algebra via a type class:

```
class Composable a where
  concatenation :: a → a → a  -- (◊), unit: empty
  conjunction  :: a → a → a  -- (∧), unit: universe
  empty        :: a           -- identity for ◊ (varepsilon)
  universe     :: a           -- identity for ∧ (neg-emptyset)
  subtraction  :: a → a → a  -- (\\)
```

concatenation sequences effects; conjunction combines simultaneous constraints; subtraction computes the residual obligation (the Brzozowski quotient for the RE instance). The identities empty and universe appear in pure.

3.2 The Pledge Type

The future field is *data-dependent*: it is a function of the return value a , allowing obligations to refer to the specific resource handle that an operation produces. The helper `evalFuture` evaluates this function at the computation's own return value.

```
data Pledge eff a = Pledge
  { ret      :: a
  , pre     :: eff      -- what must have occurred before
  , post    :: eff      -- what this computation produces
  , future  :: a → eff  -- obligation indexed by the return value
  }
```

```
evalFuture :: Pledge eff a → eff
evalFuture e = future e (ret e)
```

3.3 Monad Instances

The pure function introduces a value with no effect and no obligation:

```
pure x = Pledge { ret = x, pre = universe, post = empty, future = \_ → universe }
```

The bind operator sequences two computations; write $e_2 = f(\text{ret}(e))$.

```
e >>= f = let fe = f (ret e) in Pledge
{ ret      = ret fe
, pre      = pre e ∧ (post e \ post fe)
, post     = post e ∨ post fe
, future   = \_ → (post fe \ future e (ret e)) ∧ future fe (ret fe)
}
```

Future condition. Each future function is applied to the corresponding `ret` before the obligation is propagated:

$$\text{future}(e \gg= f) = (\text{post}(e_2) \setminus \text{future}(e)(\text{ret}(e))) \wedge \text{future}(e_2)(\text{ret}(e_2)) \quad (3)$$

This removes from e 's future obligation (evaluated at e 's return value) those traces that e_2 's postcondition satisfies, then intersects the residual with e_2 's own future obligation (evaluated at e_2 's return value). When the result normalises to $\neg\emptyset$, all obligations are discharged.

Precondition.

$$\text{pre}(e \gg= f) = \text{pre}(e) \wedge (\text{post}(e) \setminus \text{pre}(e_2)) \quad (4)$$

The term $\text{post}(e) \setminus \text{pre}(e_2)$ is the *residual precondition*: the part of e_2 's precondition not already discharged by e 's postcondition. Both must hold simultaneously in the same input history.

When $\text{post}(e)$ does not cover $\text{pre}(e_2)$, the residual is $\emptyset$, flagging a violation. When it fully covers $\text{pre}(e_2)$, the residual is ε and the combined precondition reduces to $\text{pre}(e)$ (the standard Hoare sequencing rule).

3.4 The Applicative Instance

The Pledge monad also satisfies the Applicative class, which provides `<*>` for combining a function-valued computation with an argument-valued computation *without* data-dependency between them. Writing e_f for the function computation and e_x for the argument computation:

```
ef <*> ex = Pledge
{ ret      = ret ef (ret ex)
, pre      = pre ef ∧ (post ef \ post ex)
, post     = post ef ∨ post ex
, future   = \_ → (post ex \ future ef (ret ef)) ∧ future ex (ret ex)
}
```

The structure mirrors `>>=`: preconditions and postconditions compose in the same way, and the future condition is the conjunction of the two evaluated futures after subtracting the continuation's postcondition from the first computation's future. The key difference is that `<*>` cannot make e_x *depend* on the return value of e_f ; `>>=` passes $\text{ret}(e)$ into f , enabling the data-dependent future conditions exemplified by `tlsHandshake` and `takeMVar`.

3.5 Intuition via an Example

The `takeMVar` operation returns the handle of the taken `MVar` as its return value; the future condition is data-dependent, referring to whatever handle was actually returned. For `putMVar`, the precondition requires that `takeMVar(v)` has occurred previously in the trace, and the future asserts that `putMVar(v)` must not recur until `takeMVar(v)` is called again, preventing double-put.

```
type MVar = Int
```

```
takeMVar :: MVar → Pledge RE MVar
takeMVar v = Pledge
  { ret = v, pre = universe
  , post  = Single (Atom "takeMVar" [Num v])
  , future = \r → finally (Atom "putMVar" [Num r]) } }
```

```
putMVar :: MVar → Pledge RE ()
putMVar v = Pledge
  { ret = ()
  , pre  = previously (Atom "takeMVar" [Num v])
  , post  = Single (Atom "putMVar" [Num v])
  , future = \_ → noUntil (Atom "putMVar" [Num v]) (Atom "takeMVar" [Num v]) } }
```

Good program: `v <- takeMVar 1; putMVar v`. After `takeMVar 1` ($ret = 1$) the annotations are:

$$post = takeMVar(1), \quad future = finally(putMVar(1)), \quad pre = \neg\emptyset.$$

Binding `putMVar 1` (whose $pre = previously(takeMVar(1))$, $post = putMVar(1)$):

- *Precondition.*

$$post(takeMVar(1)) \setminus pre(putMVar(1)) = takeMVar(1) \setminus previously(takeMVar(1)) = \neg\emptyset,$$

$$\text{so } pre = \neg\emptyset \wedge \neg\emptyset = \neg\emptyset. \checkmark$$

- *Future.* $putMVar(1) \setminus finally(putMVar(1)) = \neg\emptyset$; intersecting with $future(putMVar(1)) = noUntil(putMVar(1), takeMVar(1))$ gives $noUntil(putMVar(1), takeMVar(1))$. $\checkmark$
(Nullable: the program may end here; no double-put obligation is pending.)

Bad program: `putMVar 1 alone`. `pure ()` has $post = \epsilon$, so

$$\epsilon \setminus previously(takeMVar(1)) = previously(takeMVar(1)),$$

which is not nullable (the empty trace contains no `takeMVar` event). Hence

$$pre = \neg\emptyset \wedge previously(takeMVar(1)) = previously(takeMVar(1)),$$

which is not $\neg\emptyset$, so the precondition is violated. $\times$

Bad program: `missingPut (takeMVar 1 >> takeMVar 2 >> putMVar 1)`. The future of `takeMVar 2` is $finally(putMVar(2))$, which is not discharged by `putMVar 1`. The residual future retains $finally(putMVar(2))$, naming the blocked `MVar` precisely. $\times$

Bad program: `doublePutImmediate (v <- takeMVar 1; putMVar v; putMVar v)`. After the first `putMVar 1`, the future is

$$noUntil(putMVar(1), takeMVar(1)).$$

The second `putMVar 1` has $post = putMVar(1)$; subtracting it from the formula above yields $\emptyset$, so the future collapses to $\emptyset$. $\times$

4 The RE Instance

The Composable RE instance maps the algebra operations onto RE constructors and implements subtraction via derivatives:

```
instance Composable RE where
  concatenation = Seq
  conjunction  = And
  empty        = Epsilon
  universe     = anything      -- = Not Bot

  subtraction Epsilon r2 = r2
  subtraction (Not Bot) _ = Epsilon  -- Sigma* \ r = eps (trivially satisfied)
  subtraction r1 r2 =
    let alph  = atoms r1 `union` atoms r2  -- combined alphabet
        evts  = firstWith alph r1
        step e = subtraction (normalize (derivative e r1))
                              (normalize (derivative e r2))
    in foldr Or Bot (map step evts)

  normalize :: RE → RE  -- RE-specific; not part of Composable
  normalize r = {- ... algebraic simplification ... -}
```

The second clause handles the universal case: $\neg \emptyset \setminus r = \varepsilon$ for any r . The general clause collects the alphabet from *both* operands before calling `firstWith`, ensuring events appearing only in r_2 are still considered under complement via Equation (2).

The `normalize` function is defined outside the instance and applies the following RE-specific laws, which exploit the complement constructor:

Law	Applied to
$\varepsilon \cdot r = r, r \cdot \varepsilon = r$	Seq
$\emptyset \vee r = r, r \vee \neg \emptyset = \neg \emptyset$	Or
$\emptyset \wedge r = \emptyset, r \wedge \neg \emptyset = r$	And
$\neg \neg r = r$	Not
$\neg(r_1 \vee r_2) = \neg r_1 \wedge \neg r_2$	Not
$\neg(r_1 \wedge r_2) = \neg r_1 \vee \neg r_2$	Not
$\neg \emptyset = \neg \emptyset, \neg \neg \emptyset = \emptyset$	Not
$\emptyset^* = \varepsilon, \varepsilon^* = \varepsilon$	Star

5 The GuardedRE Instance

The RE instance reasons about event traces but is entirely silent about arithmetic state. The GRE instance closes this gap: a guarded regular expression pairs a trace RE with a Presburger arithmetic predicate over a concrete heap, so that a single annotation can simultaneously enforce a protocol ordering *and* an arithmetic bound.

5.1 The GuardedRE Type

```
data GuardedRE = GuardedRE PPred RE
  deriving (Eq)
```

A state (heap, trace) satisfies **GuardedRE** p r iff

$$\text{heap} \models p \quad \wedge \quad \text{trace} \in L(r).$$

The two sides are independent: p is a static constraint on the heap, while r advances through Brzozowski derivatives as events arrive.

Two smart constructors lift a pure RE or a pure Presburger predicate:

```

fromRE    :: RE    → GuardedRE
fromRE r   = GuardedRE PTrue r      -- no heap constraint

fromPPred :: PPred → GuardedRE
fromPPred p = GuardedRE p universe  -- accept any trace

```

5.2 The Composable **GuardedRE** Instance

```

instance Composable GuardedRE where
  concatenation (GuardedRE p1 r1) (GuardedRE p2 r2) =
    GuardedRE (normalizePPred (PAnd p1 p2)) (normalize (Seq r1 r2))
  conjunction (GuardedRE p1 r1) (GuardedRE p2 r2) =
    GuardedRE (normalizePPred (PAnd p1 p2)) (normalize (And r1 r2))
  empty       = GuardedRE PTrue Epsilon
  universe    = GuardedRE PTrue (Not Bot)  -- = anything
  subtraction (GuardedRE p1 r1) (GuardedRE p2 r2) =
    GuardedRE (normalizePPred (PAnd p1 p2))
      (normalize (reSubtraction r1 r2))

```

Both concatenation and subtraction conjoin both heap predicates: the residual must respect the constraints from both operands. The RE component uses the same reSubtraction and normalize as the plain RE instance.

5.3 Presburger Normalisation

Composing multiple **GuardedRE** values accumulates a tree of PAnd nodes. The normalizePPred function flattens this tree, removes duplicates, substitutes equalities (e.g. $h[a] = k \Rightarrow$ replace ValAt a by k throughout), evaluates literal comparisons, and replaces $\neg \text{true}$ with the canonical **false** constructor **PFalse**:

```

data PPred
  = PTrue | PFalse      -- canonical true/false
  | PLt PExpr PExpr | ...
  | PNot PPred | PAnd PPred PPred

normalizePPred :: PPred → PPred
normalizePPred p = rebuildAnd (substEqs eqs atoms)
where
  atoms = flattenAnd (normStep p)
  eqs   = [ (a, k) | PEq (ValAt a) (Lit k) <- atoms ]

```

The introduction of **PFalse** prevents the solver from being invoked with expressions such as $(\neg \text{true}) \wedge \dots$ that are syntactically unsatisfiable; these normalise immediately to **PFalse** without any solver call.

5.4 Example: Bounded Counter

A bounded counter stored at heap address a with maximum value M must follow the protocol $\text{init} \cdot (\text{inc} \mid \text{dec})^* \cdot \text{snapshot}$. Increment is additionally guarded by $h[a] < M$ (no overflow); decrement by $h[a] > 0$ (no underflow).

```

initCounter :: Addr → Pledge GuardedRE ()
initCounter a = Pledge
  { ret = ()
  , pre = fromRE universe
  , post = GuardedRE (PEq (ValAt a) (Lit 0))
                (Single (Atom "init" [Num a]))
  , future = \_ → fromRE universe }

increment :: Addr → Int → Pledge GuardedRE ()
increment a maxVal = Pledge
  { ret = ()
  , pre = GuardedRE (PLt (ValAt a) (Lit maxVal)) universe
  , post = fromRE (Single (Atom "inc" [Num a]))
  , future = \_ → fromPPred (PGe (ValAt a) (Lit 0)) }

```

The pre of increment carries $h[a] < M$ (pure Presburger) conjoined with the universal RE. After composing eleven increments against a bound of 10, `normalizePPred` derives the contradiction $h[a] < 10 \wedge \neg(h[a] < 10) \equiv \text{false}$, flagging the overflow before any execution takes place.

Bad programs and residuals.

Bad program	Residual pre (Presburger)	Residual pre (RE)
overflow (11 incs, max=10)	false	$\neg\emptyset$
underflow (dec at zero)	false	$\neg\emptyset$
noInit (skip init)	$\neg\emptyset$	$\emptyset$

6 The WeightedRE Instance

The RE and GRE instances are Boolean: a future condition is either satisfied or violated. The WRE instance generalises membership to a *weight* drawn from a semiring, enabling probabilistic and quantitative obligations to be expressed within the same Composable framework.

6.1 The Semiring Class

```

class (Eq w, Show w) ⇒ Semiring w where
  szero :: w          -- additive identity (blocked path)
  sone  :: w          -- multiplicative identity
  sadd  :: w → w → w
  smul  :: w → w → w

```

Setting $w = \text{Bool}$ (with $\text{sadd} = (\vee)$, $\text{smul} = (\wedge)$) recovers the Boolean RE instance. Two further instances capture quantitative scenarios:

Instance	szero	sadd	smul
Prob (probability)	0	$\min(p_1 + p_2, 1)$	$p_1 \times p_2$
Tropical (min-cost)	$+\infty$	min	+

In the probability semiring, `sadd` is bounded addition (sum of probabilities capped at 1) and `smul` is multiplication. In the tropical semiring, `sadd` is minimum (taking the cheaper alternative) and `smul` is addition (summing costs along a path).

6.2 The WRE Type and Operators

```
data WRE w
  = WBot                -- weight szero: no accepting path
  | WEps w              -- weight w: accept the empty trace
  | WSingle w Event     -- weight w: accept exactly this event
  | WSeq (WRE w) (WRE w)
  | WAdd (WRE w) (WRE w)
  | WAnd (WRE w) (WRE w)
  | WStar (WRE w)
```

The generalised nullable function accumulates the weight of all accepting paths through the empty trace:

```
wNullable :: Semiring w => WRE w -> w
wNullable (WEps w)      = w
wNullable (WSeq r1 r2)  = smul (wNullable r1) (wNullable r2)
wNullable (WAdd r1 r2)  = sadd (wNullable r1) (wNullable r2)
wNullable (WAnd r1 r2)  = smul (wNullable r1) (wNullable r2)
wNullable (WStar r)     = sone
wNullable _             = szero
```

The functions `wDerivative` and `wSubtraction` follow the same structure as their Boolean counterparts, with semiring operations substituted for Boolean ones. The `Composable (WRE w)` instance (enabled by `FlexibleInstances`) maps concatenation to `WSeq`, conjunction to `WAnd`, and subtraction to `wSubtraction`.

6.3 Derived Forms

```
wTop      :: Semiring w => WRE w
wTop      = WStar (WSingle sone Wildcard)  -- sone-weighted anything

wFinally :: Semiring w => w -> Event -> WRE w
wFinally weight e = WSeq wTop (WSingle weight e)

wGlobally :: Semiring w => w -> Event -> WRE w
wGlobally weight e = WStar (WSingle weight e)
```

6.4 Example: Probabilistic Memory (Prob Semiring)

In a probabilistic setting, each `free(a)` has an associated probability of being called. Each `malloc` call carries a future condition asserting that `free(a)` must be called with probability at least p :

```
malloc :: Addr -> Double -> Pledge (WRE Prob) Addr
malloc a p = Pledge
  { ret = a, pre = wTop
  , post  = WSingle sone (Atom "malloc" [Num a])
  , future = \r -> wFinally (Prob p) (Atom "free" [Num r]) }
```

The residual weight of `wNullable` at the end of a program measures how much of the obligation has been satisfied: a weight of `Prob 1.0` indicates that the obligation is fully discharged; a weight strictly less than 1 quantifies the probability shortfall.

6.5 Example: Min-Cost Scheduling (Tropical Semiring)

Under the tropical semiring, `WSingle (Tropical c)` `e` assigns cost `c` to emitting event `e`; `wNullable` then returns the minimum total cost of reaching an accepting state:

```
submitJob :: Int → Double → Pledge (WRE Tropical) ()
submitJob jobId cost = Pledge
  { ret = ()
  , pre = wTop
  , post = WSingle (Tropical cost) (Atom "submit" [Num jobId])
  , future = \_ →
      wFinally (Tropical 0) (Atom "complete" [Num jobId]) }
```

Composing several `submitJob` calls yields a future whose `wNullable` is the minimum total cost of completing all submitted jobs. A program that omits a complete event for some job `i` retains `wFinally (Tropical 0) (complete i)` in the residual, and the overall weight is $+\infty$ (i.e. `szero`), flagging the missing completion.

7 The SL Instance

The Composable typeclass is not inherently tied to trace-based specifications. We demonstrate its generality with a fourth instance over *separation-logic heap predicates*, where future conditions describe what heap state must hold upon completion of the computation.

7.1 Structural Parallel with RE

Table 3 shows the role each SL construct plays in the Composable interface.

Table 3. Structural parallel between the four Composable instances.

Operation	RE	GRE	WRE _w	SL
concatenation	$r_1 \cdot r_2$	$(p_1 \wedge p_2, r_1 \cdot r_2)$	<code>WSeq</code>	$P * Q$
conjunction	$r_1 \wedge r_2$	$(p_1 \wedge p_2, r_1 \wedge r_2)$	<code>WAnd</code>	$P \wedge Q$
empty	ε	$(\text{true}, \varepsilon)$	<code>WEps sone</code>	emp
universe	$\neg \emptyset$	$(\text{true}, \neg \emptyset)$	<code>wTop</code>	T
subtraction	Brzowski quotient	quotient $(\text{RE}) + \wedge (\text{PPred})$	<code>wSubtraction</code>	$P \multimap Q$

The bind formula is unchanged: $\text{pre}(e \gg f) = \text{pre}(e) \wedge (\text{post}(e) \multimap \text{pre}(e_2))$ and $\text{future}(e \gg f) = (\text{post}(e_2) \multimap \text{future}(e)) \wedge \text{future}(e_2)$. Here $\multimap$ is the magic wand (separating implication): $P \multimap Q$ holds of heap h iff for every heap h' disjoint from h satisfying P , the combined heap $h \cup h'$ satisfies Q .

7.2 The SL Predicate Type

```
data SL
= Emp          -- empty heap          (identity for SepStar)
| Top          -- any heap            (identity for Conj)
| Pure PPred   -- pure Presburger constraint (heap-independent)
| Cell Addr Val -- singleton: address a holds value v
```

```

| SepStar SL SL    -- P * Q  separating conjunction
| Conj SL SL      -- P ∧ Q  ordinary conjunction
| Wand SL SL      -- P -* Q  magic wand (residual)
deriving (Eq, Show)

```

7.3 Separating vs. Ordinary Conjunction

SepStar $P \ Q$ holds of heap h iff h splits into disjoint parts $h_1 \uplus h_2$ with $P(h_1)$ and $Q(h_2)$; it is the standard separating conjunction of ownership. Conj $P \ Q$ holds of h iff the *same* heap satisfies both P and Q . Without pure constraints, Conj on purely spatial predicates is either trivial ($\top \wedge P = P$) or unsatisfiable ($\text{Cell } l_1 \ v_1 \wedge \text{Cell } l_2 \ v_2 = \text{false}$ for $l_1 \neq l_2$). The constructor becomes useful when one side is a pure constraint $[\phi] \wedge (\ell \mapsto v)$, combining a heap-independent condition with spatial ownership.

7.4 Presburger Arithmetic for Pure Constraints

Pure constraints reuse the same PPred language introduced in Section 5 (the PExpr/PPred types and normalizePPred). The predicate Pure p holds of any heap for which evalPPred $h \ p$ is true; it enforces no ownership of cells. The smart constructor pureSL realises the standard SL box notation: pureSL **PTrue** = Emp and pureSL $p = \text{Conj } (\text{Pure } p) \ \text{Emp}$ for non-trivial p .

7.5 The Composable SL Instance

```

instance Composable SL where
  concatenation = SepStar
  conjunction   = Conj
  empty         = Emp
  universe      = Top

  subtraction Emp q = q          -- emp -* Q = Q
  subtraction Top _ = Emp        -- Top trivially covers any pre
  subtraction p   q = Wand p q  -- general: symbolic wand

```

The two base cases mirror those of the RE instance: $\varepsilon \setminus r = r$ becomes **emp** $\rightarrow^* Q = Q$ (providing no heap leaves the obligation unchanged), and $\neg \emptyset \setminus r = \varepsilon$ becomes $\top \rightarrow^* Q = \text{emp}$ (a $\top$ postcondition covers any precondition, leaving no residual). The general case records the wand symbolically for later evaluation or normalisation.

7.6 Normalisation

The normalizeSL function applies the standard SL algebraic identities:

Law	Constructor
emp $* Q = Q$, $P * \text{emp} = P$	SepStar
$\top \wedge Q = Q$, $P \wedge \top = P$	Conj
$P \wedge P = P$ (idempotent)	Conj
emp $\rightarrow^* Q = Q$	Wand
$P \rightarrow^* \top = \top$	Wand

All rules apply recursively bottom-up.

8 Examples

We demonstrate Pledge across ten use cases spanning all four instances. Each example defines primitive operations annotated with pre, post, and future, composes them in the Pledge monad, and inspects the normalised annotations of the result. A future of $\neg\emptyset$ (or `wNullable = some` for WRE) and a pre of $\neg\emptyset$ together certify temporal correctness.

8.1 TLS Handshake Lifecycle

The `tls` library obligation is representative of a broad class of Haskell protocol obligations: a successful handshake must always be paired with a subsequent bye. We model the Context as an integer identifier.

```
type Context = Int

tlsHandshake :: Context → Pledge RE Context
tlsHandshake ctx = Pledge
  { ret    = ctx
  , pre    = universe
  , post   = Single (Atom "handshake" [Num ctx])
  , future = \c → finally (Atom "bye" [Num c]) }

tlsBye :: Context → Pledge RE ()
tlsBye ctx = Pledge
  { ret    = ()
  , pre    = Single (Atom "handshake" [Num ctx])
  , post   = Single (Atom "bye" [Num ctx])
  , future = \_ → universe }

tlsSend :: Context → Pledge RE ()
tlsSend ctx = Pledge
  { ret    = ()
  , pre    = Single (Atom "handshake" [Num ctx])
  , post   = Single (Atom "send" [Num ctx])
  , future = \_ → universe }

-- Good: handshake followed by data exchange then bye
goodSession = do
  ctx <- tlsHandshake 1
  tlsSend ctx
  tlsBye ctx

-- Bad future: bye omitted; residual = finally(bye(1))
missingBye = do
  ctx <- tlsHandshake 1
  tlsSend ctx

-- Bad future: two contexts opened, only one closed
partialClose = do
  ctx1 <- tlsHandshake 1
```

```
ctx2 <- tlsHandshake 2
tlsBye ctx1
```

The program `missingBye` retains `finally(bye(1))` as its residual future, naming the unclosed context precisely. Meanwhile, `partialClose` retains `finally(bye(2))`: the future condition is data-dependent, so the residual identifies context 2 specifically, not a generic “some context was not closed.”

8.2 MVar Lifecycle

The `takeMVar` operation returns the taken handle; its data-dependent future condition uses the lambda `r` so the obligation tracks the *returned* handle. The `putMVar` operation carries as its future obligation `noUntil(putMVar(v), takeMVar(v))`, forbidding a second `putMVar(v)` until `takeMVar(v)` resets the guard. (The full definitions of `takeMVar` and `putMVar` appear in Section 3.2.)

```
-- Good: data-dependent - obligation discharged on the returned handle
takeAndPut n = do { v <- takeMVar n; putMVar v }

-- Good: re-acquire same MVar after releasing - noUntil is reset by takeMVar
takePutTakePut = do { v <- takeMVar 1; putMVar v; v <- takeMVar 1; putMVar v }

-- Bad future: putMVar(2) obligation remains (MVar 2 never returned)
missingPut = do { v1 <- takeMVar 1; putMVar v1; _ <- takeMVar 2; return () }

-- Bad pre: takeMVar(1) precondition not met
putWithoutTake = putMVar 1

-- Bad future: double-put immediately collapses future to Bot
doublePutImmediate = do { v <- takeMVar 1; putMVar v; putMVar v }

-- Bad future: double-put with intervening work, still collapses to Bot
doublePutWithWork = do
  { v1 <- takeMVar 1; v2 <- takeMVar 2; putMVar v1; putMVar v2; putMVar v1 }
```

The three examples below return non-unit values, showing that future conditions are tracked regardless of the return type.

```
-- Good: ret = MVar; future = finally(putMVar(1)) - obligation visible in residual
takeReturnHandle :: Pledge RE MVar
takeReturnHandle = takeMVar 1

-- Bad future: ret = (MVar,MVar); future = finally(putMVar(1)) ^ finally(putMVar(2))
takeTwoReturnPair :: Pledge RE (MVar, MVar)
takeTwoReturnPair = do { v1 <- takeMVar 1; v2 <- takeMVar 2; return (v1, v2) }

-- Good: ret = MVar; future = noUntil(putMVar(1),takeMVar(1)) - no pending finally
takePutReturnHandle :: Pledge RE MVar
takePutReturnHandle = do { v <- takeMVar 1; putMVar v; return v }
```

8.3 Mutex / Lock Lifecycle


```
release :: Int → Pledge RE ()
release m = Pledge
  { ret = ()
  , pre  = Single (Atom "acquire" [Num m])
  , post = Single (Atom "release" [Num m])
  , future = universe }

-- Bad future: lock 1 not released
lockLeak = do { acquire 1; acquire 2; release 2 }

-- Bad pre: release without acquire
releaseWithoutAcquire = release 1
```

8.4 Database Transactions

The `beginTx` operation carries a *disjunctive* future condition requiring that the transaction is eventually either committed or rolled back; `commit` and `rollback` each discharge this obligation:

```
beginTx :: Pledge RE ()
beginTx = Pledge
  { ret = (), pre = universe
  , post  = Single (Atom "beginTx" [])
  , future = Or (finally (Atom "commit" []))
                (finally (Atom "rollback" [])) }

commit :: Pledge RE ()
commit = Pledge
  { ret = (), pre = Single (Atom "beginTx" [])
  , post  = Single (Atom "commit" [])
  , future = universe }

rollback :: Pledge RE ()
rollback = Pledge
  { ret = (), pre = Single (Atom "beginTx" [])
  , post  = Single (Atom "rollback" [])
  , future = universe }

-- Bad future: transaction never closed
openTx = do { beginTx }

-- Good: transaction properly committed
committedTx = do { beginTx; commit }

-- Good: transaction properly rolled back
abortedTx = do { beginTx; rollback }
```

`openTx` retains the full disjunction $\text{finally}(\text{commit}) \vee \text{finally}(\text{rollback})$ as its residual future. `committedTx` and `abortedTx` fully discharge the obligation, yielding $\neg\emptyset$.

RE summary. Table 4 summarises the three RE domains.

Table 4. RE instance: bad programs and residual annotations.

Domain	Bad program	Residual future	Residual pre
Memory	missingFree	finally(free(2))	$\neg\emptyset$
Memory	freeWithoutMalloc	$\neg\emptyset$	$\emptyset$
Memory	doubleFreeImmediate	$\emptyset$	$\neg\emptyset$
Memory	leakAndDoubleFree	$\emptyset$	$\neg\emptyset$
Memory	allocReturnAddr	finally(free(1))	$\neg\emptyset$
Memory	allocTwoReturnPair	finally(free(1)) $\wedge$ finally(free(2))	$\neg\emptyset$
Mutex	acquire 1 leaked	finally(release(1))	$\neg\emptyset$
Mutex	release without acquire	$\neg\emptyset$	$\emptyset$
Transaction	no commit/rollback	finally(commit) $\vee$ finally(rollback)	$\neg\emptyset$

8.5 Bounded Counter (GRE Instance)

The bounded counter (§5) demonstrates that both the Presburger and RE components are necessary. The RE alone would accept overflow sequences; the Presburger predicate alone would accept arbitrarily ordered events.

```
-- Good: init, two increments, one decrement, snapshot
normalUse :: Pledge GuardedRE ()
normalUse = do
  initCounter 0; increment 0 10; increment 0 10
  decrement 0; snapshot 0

-- Bad: 11 increments against max 10
overflow :: Pledge GuardedRE ()
overflow = do
  initCounter 0
  sequence_ (replicate 11 (increment 0 10))
  snapshot 0
```

For overflow, the combined pre after composing all eleven increments normalises to $[\text{false}] \wedge \neg\emptyset$: the Presburger side collapses to **PFalse** while the RE side remains the universal language, pinpointing the arithmetic violation.

GRE *summary*.

Bad program	Residual pre (GRE)	Cause
overflow	$[\text{false}] \neg\emptyset$	$h[0] < 10$ violated
underflow	$[\text{false}] \neg\emptyset$	$h[0] > 0$ violated
noInit	$[\text{true}] \emptyset$	init event missing

8.6 Task Scheduling (WRE Instance)

The task-scheduling example shows that the tropical semiring naturally accumulates minimum costs. Each submitJob operation carries both a submission cost and a future condition requiring that every submitted job eventually completes.

```
-- Good: submit two jobs, complete both
goodSchedule :: Pledge (WRE Tropical) ()
goodSchedule = do
```

```

submitJob 1 2.0; submitJob 2 3.0
completeJob 1;   completeJob 2

-- Bad: job 2 never completed
missingComplete :: Pledge (WRE Tropical) ()
missingComplete = do
  submitJob 1 2.0; submitJob 2 3.0
  completeJob 1

```

For `goodSchedule`, `wNullable (evalFuture ...)` returns **Tropical** 0, indicating all obligations are discharged at zero residual cost. For `missingComplete`, the residual future still contains `wFinally (Tropical 0) (complete 2)`, and the overall weight is $+\infty$, flagging job 2 as unfinished.

The probability-semiring memory example follows analogously: a call `malloc a p` with $p = 0.9$ leaves a residual weight of 0.9 if `free a` is never called, dropping to 0 when it is.

8.7 Heap Memory (SL Instance)

The heap-memory example mirrors the RE memory example but uses heap ownership rather than event traces.

```

alloc :: Addr → Val → Pledge SL ()
alloc a v = Pledge
  { ret = (), pre = Top, post = Cell a v, future = Top }

free :: Addr → Val → Pledge SL ()
free a v = Pledge
  { ret = (), pre = Cell a v, post = Emp, future = Top }

writeCell :: Addr → Val → Val → Pledge SL ()
writeCell a old new = Pledge
  { ret = (), pre = Cell a old, post = Cell a new, future = Top }

```

Good: `alloc 0 1` $\gg$ `alloc 1 2`. The combined post is `SepStar (Cell 0 1) (Cell 1 2)`: two disjoint ownerships composed by separating conjunction, correctly reflecting that the two cells reside at distinct addresses.

Bad: `leak (alloc 0 1` $\gg$ `alloc 1 2` $\gg$ `free 0 1)`. After freeing only address 0, the post retains `Cell 1 2`, naming the unreleased ownership.

Bad: `read without ownership (readCell 0 42)`. The pre of the bare read is `Cell 0 42`, which is never discharged, flagging the missing ownership.

8.8 Bank Account with Presburger Guards (SL Instance)

The bank account example demonstrates Pure Presburger constraints inside `Conj`.

```

withdraw :: Addr → Val → Val → Pledge SL ()
withdraw a old amount = Pledge
  { ret = ()
  , pre = Conj (Pure (PGe (ValAt a) (Lit amount))) (Cell a old)
  , post = Cell a (old - amount)
  , future = Top }

closeAccount :: Addr → Val → Pledge SL ()

```

```

closeAccount a bal = Pledge
{ ret = ()
, pre = Conj (Pure (PEq (ValAt a) (Lit 0))) (Cell a bal)
, post = Emp, future = Top }

```

The precondition of `withdraw` is a `Conj` of a pure Presburger guard ($h[a] \geq \text{amount}$) and a spatial ownership claim (`Cell a old`). Without the `Pure` constructor, `Conj` of two purely spatial predicates would be either trivial or unsatisfiable (see §7); the combination of `Pure` and `Conj` is what makes value-dependent preconditions expressible.

Bad: overdraft. After `deposit 0 0 100`, the account holds 100. The call `withdraw 0 100 150` carries $pre = \text{Conj} (\text{Pure} (h[0] \geq 150)) (\text{Cell } 0 \ 100)$. Since $100 < 150$, the pure guard is not satisfiable, and the residual pre of the composed program retains this unsatisfiable constraint.

8.9 Linked List (SL Instance)

A singly-linked-list node occupies two consecutive cells: address a holds the payload value and address $a+1$ holds the next pointer.

```

allocNode :: Addr → Val → Val → Pledge SL ()
allocNode a val next = Pledge
{ ret = (), pre = Top
, post = SepStar (Cell a val) (Cell (a+1) next)
, future = Top }

```

Two-node list (`allocNode 0 10 2 >> allocNode 2 20 (-1)`). The combined post is a four-cell `SepStar`:

$$(\text{Cell } 0 \ 10 * \text{Cell } 1 \ 2) * (\text{Cell } 2 \ 20 * \text{Cell } 3 \ -1)$$

This is the canonical representation of a two-node list in separation logic: all four cells are disjointly owned, and the structure of `SepStar` mirrors the structure of the list.

SL summary. Table 5 summarises the three SL domains. The key difference from the RE instance is that residual annotations are heap predicates rather than traces: a leaked cell appears in post, an unsatisfiable ownership constraint in pre.

Table 5. SL instance: bad programs and residual annotations.

Domain	Bad program	Residual post	Residual pre
Heap memory	<code>Cell 1 2</code> not freed	<code>Cell 1 2</code> (leaked)	$\top$
Heap memory	read without ownership	emp	<code>Cell 0 42</code>
Bank account	insufficient balance	<code>Cell 0 -50</code>	unsatisfiable guard
Bank account	close with non-zero balance	emp	unsatisfied $h[0] = 0$
Linked list	read without ownership	emp	<code>Cell 0 99</code>

9 Monad Laws

Every monad must satisfy three equational laws that together guarantee sequential composition behaves predictably:

- **Left identity:** $\text{pure } x \gg= f = f \ x$. Wrapping a value in `pure` and immediately binding it is the same as applying f directly; `pure` introduces no effect of its own.

- **Right identity:** $e \gg \text{pure} = e$. Binding a computation to `pure` is the same as the computation itself; `pure` consumes no observable state.
- **Associativity:** $(e \gg f) \gg g = e \gg (\lambda x \rightarrow f\ x \gg g)$. Re-parenthesising a chain of binds does not change its meaning; sequential composition is order-preserving but grouping-independent.

When these laws hold, programmers can freely refactor monadic chains without changing the observable semantics: extracting sub-expressions, fusing steps, or reordering `do`-blocks are all valid.

What the laws check. Because `future` has type $a \rightarrow \text{eff}$, the laws are stated in terms of `evalFuture` $e = \text{future } e \text{ (ret } e)$, written $\widehat{\text{future}}(e)$ for brevity. The laws hold for the triple $(\text{ret}, \text{post}, \widehat{\text{future}})$ assuming `Composable` forms a monoid under $\diamond$ with unit `empty`, and `universe` is the unit for $(\wedge \setminus)$. The pre field satisfies the Hoare-rule property (Equation 4) rather than strict monad-law equality, because it accumulates residual precondition obligations; we return to this at the end of the section.

Left identity. $\text{pure } x \gg f = f\ x$.

ret: $\text{ret}(f\ x)$. $\checkmark$ *post:* $\varepsilon <> \text{post}(f\ x) = \text{post}(f\ x)$. $\checkmark$

For the future, let $fe = f(\text{ret}(\text{pure } x)) = f\ x$. Since $\widehat{\text{future}}(\text{pure } x) = \neg\emptyset$:

$$\begin{aligned} \widehat{\text{future}}(\text{pure } x \gg f) &= (\text{post}(fe) \setminus \widehat{\text{future}}(\text{pure } x)) \wedge \widehat{\text{future}}(fe) \\ &= (\text{post}(f\ x) \setminus \neg\emptyset) \wedge \widehat{\text{future}}(f\ x) = \varepsilon \wedge \widehat{\text{future}}(f\ x) = \widehat{\text{future}}(f\ x). \checkmark \end{aligned}$$

Right identity. $e \gg \text{pure} = e$.

ret: $\text{ret}(e)$. $\checkmark$ *post:* $\text{post}(e) <> \varepsilon = \text{post}(e)$. $\checkmark$

Let $fe = \text{pure}(\text{ret}(e))$, so $\text{post}(fe) = \varepsilon$ and $\widehat{\text{future}}(fe) = \neg\emptyset$. Then:

$$\widehat{\text{future}}(e \gg \text{pure}) = (\varepsilon \setminus \widehat{\text{future}}(e)) \wedge \neg\emptyset = \widehat{\text{future}}(e) \wedge \neg\emptyset = \widehat{\text{future}}(e). \checkmark$$

Associativity. Re-parenthesising $(e \gg f) \gg g$ as $e \gg (\lambda x \rightarrow f\ x \gg g)$ must not change the specification.

ret and *post:* both sides yield $\text{ret}(g\ r_2)$ and $\text{post}(e) <> \text{post}(f\ r_1) <> \text{post}(g\ r_2)$, following from associativity of $\diamond$. $\checkmark$

For the future, let $r_1 = \text{ret}(e)$, $r_2 = \text{ret}(f\ r_1)$. Unfolding the bind formula (Equation (3)) on both sides and applying associativity of $\wedge$, both reduce to the same expression:

$$(\text{post}(g\ r_2) \setminus (\text{post}(f\ r_1) \setminus \widehat{\text{future}}(e))) \wedge \widehat{\text{future}}(f\ r_1) \wedge \widehat{\text{future}}(g\ r_2). \checkmark$$

Precondition. The pre field satisfies the Hoare-rule property (Equation (4)) rather than strict monad-law equality. The residual $\text{post}(e) \setminus \text{pre}(e_2)$ captures which part of e_2 's precondition e 's postcondition has not discharged: full coverage gives the standard Hoare sequencing rule; absent coverage collapses to $\emptyset$, flagging the violation immediately.

10 Related Work

Pre/post-condition contracts in Haskell. *Liquid Haskell* [Vazou et al. 2014] expresses pre/post-conditions as refinement types verified by an SMT solver. Findler and Felleisen [Findler and Felleisen 2002] establish a runtime contract discipline for higher-order functions. QuickCheck [Claessen and Hughes 2000] supports conditional properties via the implication combinator `==>`, which restricts a property to inputs satisfying a given precondition. Swierstra and Baanen [Swierstra and Baanen 2019] embed Hoare triples as an indexed monad, propagating pre/post annotations through

monadic bind analogously to `Pledge`. `Pledge` complements these by adding a *temporal* dimension through the future condition, without requiring type-system extensions.

Future conditions and temporal verification. `ProveNFix` [Song et al. 2024] encodes temporal properties as regular expressions to guide automated program repair. `Pledge` is complementary in orientation: rather than repairing programs after the fact, it carries temporal obligations as first-class annotations so that violations are identified through annotation inspection, prior to any effectful execution. Song et al. [2025] propose future conditions as a first-class extension to pre/post-condition contracts, formalising the notion that an operation may discharge temporal obligations onto its continuation. `Pledge` realises this idea as a Haskell library, contributing a monadic interface, derivative-based propagation, and four concrete `Composable` instances that the formal system leaves unaddressed.

Parameterised and graded monads. Atkey [Atkey 2009] introduces parameterised monads, written $T(S_1, S_2, A)$, for Hoare-logic-style pre/post state tracking. Katsumata [Katsumata 2014] generalises this to graded monads over a monoid; Orchard and Wu [Orchard and Yoshida 2016] show effect systems and session types are instances. `Pledge` is related but orthogonal: rather than indexing by effect *type* at the type level, we carry an explicit obligation as `future :: a → eff` at the value level, preserving compatibility with the standard `Monad` class, `do`-notation, and the `mtl` stack.

Resource management in Haskell. The bracket combinator lexically scopes `cleanup` but cannot propagate obligations beyond the lexical scope. The `ResourceT` [Snoyman 2011] library threads a finaliser queue through every bind, discharging `cleanup` when the block exits. Both `Acquire` and `Managed` generalise this to an applicative/monadic interface. All three handle the *cleanup* half of a resource obligation but cannot express *protocol* or *quantitative* obligations, which require explicit future-condition propagation across bind steps.

Linear types. Linear Haskell [Bernardy et al. 2018] introduces multiplicity annotations enforced by GHC since version 9.0; `linear-base` [Tweag and contributors 2021] provides a `RIO` monad with linear file handles, statically guaranteeing every opened handle is consumed exactly once. Linear types enforce *structural* obligations (use exactly once), whereas future conditions enforce *temporal* obligations (use in a specified order, or with a specified probability). A hybrid combining linear types for uniqueness with `Pledge` for temporal ordering could achieve stronger guarantees than either alone.

Effect systems and algebraic effect handlers. Algebraic effect handlers [Pretnar 2015] decompose effectful computations into an operation set and a separate handler; resource lifecycle obligations are typically encoded in the handler’s state machine. `Polysemy`’s `Scoped` effect attaches resource creation and `cleanup` to a lexical scope threaded through the effect row. `Pledge` takes the complementary route: obligations are values in the annotation fields, composable across any handler.

Session types. Session types [Gay and Hole 2005; Honda 1993] enforce communication protocols at the type level. The `sessiontypes` library [van Walree 2017] realises this via an indexed monad `STTerm m s s' a`; the type checker rejects programs that violate the protocol order. `Pledge` is more lightweight: it operates within the standard `Monad` class, admits quantitative and heap-ownership obligations that session types do not, and is applicable beyond message-passing.

Typestate. Typestate systems [Bierhoff and Aldrich 2007; Strom and Yemini 1986] encode resource lifecycle protocols in the type system. Brady’s `ST` monad encoding in `Idris` [Brady 2013] demonstrates full compile-time enforcement in a dependently typed language. `Pledge` achieves comparable expressiveness at the library level using standard Haskell type classes, at the cost of moving

checking to annotation-inspection time rather than type-inference time: violations surface as non-trivial residual annotations in pre or future, which the programmer can evaluate before executing effects rather than relying on the type checker to reject programs.

Runtime monitoring. Runtime verification [Leucker and Schallhart 2009] checks temporal properties by instrumenting a concrete execution. Pledge propagates obligations in the specification layer: residuals surface as soon as annotations are composed, prior to any effectful run, without requiring a concrete execution.

Brzowski derivatives. Derivatives of regular expressions [Brzowski 1964] have found application in parsing [Might et al. 2011; Owens et al. 2009] and model checking. Equation (1) extends this to complement via $\partial_a(\neg r) = \neg(\partial_a r)$; our subtraction operator is a novel application to future-condition propagation in a monadic setting.

11 Discussion and Future Work

Decidability. Equality of normalised regular expressions is decidable, so the fixed-point computation in subtraction terminates whenever the derivative sequence is finite, which holds for the regular languages used in our examples. Extending to context-free or ω -regular specifications remains future work.

Precondition monad laws. The pre field satisfies the Hoare-rule property rather than strict monad-law equality (Section 9). A natural refinement is to make pre a separate, non-monadic annotation computed by a static analysis pass over the Pledge term, separating the behavioural monad from the precondition contract checker.

Integration with effect handlers. Lifting Pledge into an effect-handler framework (effectful or polysemy) would allow temporal specifications to compose alongside algebraic effect rows. A lightweight *shadow* approach requires no deep integration: keep the Pledge RE computation as a pure specification evaluated as obligation values before the run, while a separate Eff es computation carries out the real effects, with assertions linking the two at the top level:

```
main = do
  assert (normalize (pre      spec) == top) "precondition_violated"
  assert (normalize (evalFuture spec) == top) "future_obligation_unmet"
  runEff impl
```

Annotation checking is entirely pure; no changes to the effect stack are required.

Closed-world alphabet assumption. The firstWith function uses $\Sigma_0 = \text{atoms}(r_1) \cup \text{atoms}(r_2)$ as its working alphabet. This is sound under the *closed-world assumption* that every relevant event is mentioned in the two operands. A fully open-world implementation would accept an explicit Σ declaration, enabling complement over an unrestricted domain.

Connection to indexed and graded monads. Because `future :: a → eff` indexes the obligation by the return value, Pledge echoes the indexed/graded-monad literature where types are indexed by computational effects; here the index is a return value instead of an effect row.

Specification versus runtime enforcement. Pledge is a *specification* tool: the pre, post, and future fields are pure Haskell values that describe what a computation *should* do, not runtime guards that intercept execution. In particular, Pledge does not integrate with GHC's asynchronous-exception machinery. If an async exception fires mid-computation and the continuation never calls putMVar, Pledge's future condition correctly flags the *specification* as carrying an unmet obligation, but

no runtime mechanism prevents the violation from occurring. Pairing Pledge specifications with runtime enforcement via bracket or mask (so that obligations are both expressed and enforced) remains future work.

Solver-backed discharge of residual conditions. Residual annotations are currently surfaced as symbolic expressions. A natural extension would connect them to an automatic solver: RE residuals checked by NFA/DFA construction or a BDD library; SL predicates with Presburger guards discharged by Z3 or CVC5, reporting concrete counterexample traces or unsatisfiable heaps.

Parallel composition. The Pledge monad is inherently sequential; concurrent programs require a parallel combinator $e_1 \parallel e_2$. In the RE instance this corresponds to the shuffle product $r_1 \sqcup r_2$; the SL instance admits a parallel connective via the concurrent separation logic frame rule [O’Hearn 2007]. Extending Composable with a parallel operation is a natural next step.

Quantitative future conditions. A natural further direction is to connect WRE to probabilistic temporal logics (PCTL) or quantitative Kleene algebras; in the SL instance, a quantitative analogue could express amortised resource bounds or expected heap usage.

12 Conclusion

We have presented *Pledge*, a Haskell library that augments effectful computations with *future conditions*: annotations that constrain what the remainder of a program must accomplish.

The central insight is that trace-based, arithmetic, quantitative, and heap-ownership obligations are structurally identical at the level of a five-operation Composable algebra. All four instances (RE, GRE, WRE, SL) share a single monadic bind formula, and the diagnostic signal is uniform: a non-trivial residual in pre or future identifies exactly which obligation remains unmet and why.

Pledge requires no language extensions beyond standard type classes, composes via a standard Monad instance, and integrates transparently with the `mtl` stack. Immediate directions for future work include parallel composition (shuffle product in the RE instance; frame rule in the SL instance), solver-backed residual discharge, and integration with algebraic effect handlers.

References

- Robert Atkey. 2009. Parameterised Notions of Computation. *Journal of Functional Programming* 19, 3–4 (2009), 335–376. doi:10.1017/S095679680900728X
- Jean-Philippe Bernardy, Mathieu Boespflug, Ryan R. Newton, Simon Peyton Jones, and Arnaud Spiwack. 2018. Linear Haskell: Practical Linearity in a Higher-Order Polymorphic Language. *Proceedings of the ACM on Programming Languages* 2, POPL, Article 5 (2018), 29 pages. doi:10.1145/3158093
- Kevin Bierhoff and Jonathan Aldrich. 2007. Modular Typestate Checking of Aliased Objects. In *Proceedings of the 22nd Annual ACM SIGPLAN Conference on Object-Oriented Programming Systems and Applications (OOPSLA)*. 301–320. doi:10.1145/1297027.1297050
- Edwin Brady. 2013. Idris, a General-Purpose Dependently Typed Programming Language: Design and Implementation. *Journal of Functional Programming* 23, 5 (2013), 552–593. doi:10.1017/S095679681300018X
- Janusz A. Brzozowski. 1964. Derivatives of Regular Expressions. *J. ACM* 11, 4 (1964), 481–494. doi:10.1145/321239.321249
- Koen Claessen and John Hughes. 2000. QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs. In *Proceedings of the 5th ACM SIGPLAN International Conference on Functional Programming (ICFP)*. 268–279. doi:10.1145/351240.351266
- Robert Bruce Findler and Matthias Felleisen. 2002. Contracts for Higher-Order Functions. In *Proceedings of the 7th ACM SIGPLAN International Conference on Functional Programming (ICFP)*. 48–59. doi:10.1145/581478.581484
- Simon J. Gay and Malcolm Hole. 2005. Subtyping for Session Types in the Pi Calculus. *Acta Informatica* 42, 2-3 (2005), 191–225. doi:10.1007/s00236-005-0177-z
- Kohei Honda. 1993. Types for Dyadic Interaction. In *Proceedings of the 4th International Conference on Concurrency Theory (CONCUR) (Lecture Notes in Computer Science, Vol. 715)*. 509–523. doi:10.1007/3-540-57208-2_35
- Shin-ya Katsumata. 2014. Parametric Effect Monads and Semantics of Effect Systems. *ACM SIGPLAN Notices* 49, 1 (2014), 633–646. doi:10.1145/2578855.2535846

- Martin Leucker and Christian Schallhart. 2009. A Brief Account of Runtime Verification. *Journal of Logic and Algebraic Programming* 78, 5 (2009), 293–303. doi:10.1016/j.jlap.2008.08.004
- Matthew Might, David Darais, and Daniel Spiewak. 2011. Parsing with Derivatives: A Functional Pearl. In *Proceedings of the 16th ACM SIGPLAN International Conference on Functional Programming (ICFP)*. 189–195. doi:10.1145/2034773.2034801
- Peter W. O’Hearn. 2007. Resources, Concurrency, and Local Reasoning. *Theoretical Computer Science* 375, 1–3 (2007), 271–307. doi:10.1016/j.tcs.2006.12.035
- Dominic Orchard and Nobuko Yoshida. 2016. Effects as Sessions, Sessions as Effects. In *Proceedings of the 43rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL)*. 568–581. doi:10.1145/2837614.2837634
- Scott Owens, John Reppy, and Aaron Turon. 2009. Regular-expression Derivatives Re-examined. *Journal of Functional Programming* 19, 2 (2009), 173–190. doi:10.1017/S0956796808007090
- Matija Pretnar. 2015. An Introduction to Algebraic Effects and Handlers. Invited tutorial paper. *Electronic Notes in Theoretical Computer Science* 319 (2015), 19–35. doi:10.1016/j.entcs.2015.12.003
- Michael Snoyman. 2011. conduit: Streaming Data Processing. <https://hackage.haskell.org/package/conduit>. Haskell package on Hackage; introduces the ResourceT transformer.
- Yahui Song, Darius Foo, and Wei-Ngan Chin. 2025. Specifying and Verifying Future Conditions. In *Static Analysis — 32nd International Symposium, SAS 2025 (Lecture Notes in Computer Science, Vol. 16100)*. Springer, 90–112. doi:10.1007/978-3-032-07106-4_5
- Yahui Song, Xiang Gao, Wenhua Li, Wei-Ngan Chin, and Abhik Roychoudhury. 2024. ProveNFix: Temporal Property-Guided Program Repair. *Proceedings of the ACM on Software Engineering* 1, FSE, Article 11 (2024), 23 pages. doi:10.1145/3643737
- Robert E. Strom and Shaula Yemini. 1986. Typestate: A Programming Language Concept for Enhancing Software Reliability. *IEEE Transactions on Software Engineering* 12, 1 (1986), 157–171. doi:10.1109/TSE.1986.6312929
- Wouter Swierstra and Tim Baanen. 2019. A Predicate Transformer Semantics for Effects (Functional Pearl). *Proceedings of the ACM on Programming Languages* 3, ICFP, Article 103 (2019), 26 pages. doi:10.1145/3341707
- Tweag and contributors. 2021. linear-base: Standard Library for Linear Types in Haskell. <https://github.com/tweag/linear-base>.
- Ferdinand van Walree. 2017. sessiontypes: Session Types for Haskell. <https://hackage.haskell.org/package/sessiontypes>.
- Niki Vazou, Eric L. Seidel, Ranjit Jhala, Dimitrios Vytiniotis, and Simon Peyton Jones. 2014. Refinement Types for Haskell. In *Proceedings of the 19th ACM SIGPLAN International Conference on Functional Programming (ICFP)*. 269–282. doi:10.1145/2628136.2628161